

Cecilia

Account and payment questions,
answered with verified data

Spanish and Portuguese · Mexico, Colombia, Argentina

Team Cecilia

cecil-ai.onrender.com



Account and payment demand

35%

of 686,296 contacts

240,056 account and payment inquiries

About 411 agent hours per month

Median demand · handling time only

Measured service	Account / payment
First-contact resolution	91.5%
Average queue time	120 seconds
Average call duration	221 seconds
Customer satisfaction	2.91 / 5
Share arriving by phone	85%

Synthetic historical data. Queue times are similar across reasons; this dataset does not establish that waiting causes low satisfaction.

Model interpretation and verified actions

01 Identity and policy

Ownership checks and safety guards run in code.

02 Model chooses the lookup


Masked identifiers and product aliases. No customer records.

03 Code writes the reply

Verified tool results or fixed Spanish / Portuguese templates.

04 Customer confirms the action

A trace opens only after a plain yes, then a service read-back.



Captured on the deployed demo · 4 October 2026

One primary model response per turn; bounded retries and fallback may add attempts. No money moves.

Evaluation on the same 548 cases

Metric	Keyword bot	Sonnet 5	Haiku 4.5
Safe resolution / 238 in scope	70.2%	97.1%	79.0%
Required escalations missed / 168	72	1	37
Unsafe outcomes / 548, runs 1 / 2 / 3	0 (1 run)	0 / 0 / 0	0 / 1 / 0
Cases sending customer records / 548	Not applicable	0, every run	0, every run
Latency p50 / p95 per case	11 / 49 ms	1.4 / 3.0 s	1.1 / 4.0 s
Estimated USD per safe resolution	No model	0.0034	0.0084
Safe resolution across three runs	Deterministic	96.6–97.5%	78.6–80.7%

ES + PT, team-written test conversations. Run 1 unless stated. Sonnet 5: 231 / 238 safe resolutions, 95% CI 94.0–98.6%.

Offline adversarial model: 0 unsafe / 548. Zero observed failures does not establish zero risk.

4 Oct 2026 reports · Costs use assumed Sep 2026 list prices, not invoices. Test results are not production outcomes.

Data quality and engineering evidence

4.4M transactions - 150,000 customers

Full load: 294 checks, 0 errors, 16 warnings; 1 check not run.

Contracts and recovery

Quarantine gate, rollback, row and partition lineage; late-arrival fixture.

Model selection

Intent classifier: 84.7% vs 63.5% keywords on 85 texts. MLflow tracks prompt, model, code and data hashes.

Known data limitations

57% missing USD amounts; 42 distinct transcript texts; no Portuguese in the supplied corpus.

Classifier comparison is post-hoc: two keyword patterns touched the test split. Demo serves a 5,000-customer sample, as of 17 June 2026.

The deployed demo and the next steps

cecil-ai.onrender.com

Guided scenarios, customer chat and bank console

Verified 4 October 2026

7 scenarios · 10 turns passed

Live model: claude-sonnet-5

ES + PT · confirmed tracing · safe fallback

Before production

Bank identity provider; Redis sessions; PII protection; operational alerts; voice support.

Next workflow: transaction disputes.

Public repository: github.com/try-cecilia/factored-hackathon-2026-cecilia

Synthetic demo data and test identities. No real banking transactions. Recorded model outage is deliberately simulated.